

# Going Concurrent with Python

Mikko Harju

Taiste

October 12, 2011

# About me

- I'm Mikko Harju, Technology Director at Taiste
- Python user since 2008
- Functional programming enthusiast

# Agenda

- We'll look slightly at *threading* and more at *processes* and how do they differ when implementing concurrency in Python.

# Threads

- Threads have classically been the basic construct for parallel programming.
- They live within the process boundary, hence sharing the memory space.
- This makes it fairly trivial to transfer information from one thread to another.
- However, hilarity ensues when mutable state enters the equation...

# Threads

- To prevent problems when writing data from multiple threads to a shared state variable we need to introduce different kinds of *locking mechanisms*.
- Possibilities include e.g. mutexes, semaphores and critical sections.
- However, if the threads only access the data in read-only way, there is no need for locking.
- This is called "Shared-nothing" -approach. For example Erlang kinda has this with its "lightweight processes".

# GIL

- GIL is a mechanism to ensure that threads co-operate nicely with non-thread safe constructs
- This is done by locking down the interpreter by giving exclusive access to one thread at a time.
- This effectively means that only one CPU bound thread is actually doing anything useful in one python interpreter process at a time.
- When doing I/O, the interpreter can release the lock. The interpreter also has periodic checks to go along with this to make it possible to parallelize CPU bound threads.

# GIL

- When we add more processors (or cores) to the mix, things get more complicated.
- $N$  threads can be scheduled simultaneously on  $N$  processors, making them all compete over the GIL. Nice.
- So really, threads are not the right way to go in Python.
- There are some use cases where they might come in handy, where the problem is more I/O bound than CPU bound.
- But let us concentrate on the more fruitful (and the right way if you ask me) of doing real multiprocessor concurrency: *processes*

# Processes

- Processes are OS backed construct.
- Each process has its own memory space, stack, registers and that kind of stuff.
- Scheduling is performed by the operating system.



# Let's do this!

- Nothing beats practice, so let's do parallel image processing with PIL using the multiprocessing package and see what we come up with.
- Prepare your terminals :-)